

Cloud Computing Assignment2

All commands below based on Ubuntu 22.04 LTS and AMD64 CPU

Part 1: Creating an Application:

1. Get App from this link:

https://drive.google.com/drive/u/0/folders/19AZciSY1l2neLVfJUzqRG4a_mDstnaK_

2. Install or start MongoDB

```
# All the following commands will run at Linux
sudo systemctl start mongod
```

```
#or we can start with brew
brew services start mongodb-community@6.0
```

3. Run the application

```
python3 app.py
```

Test the Todo List APP

Part 2: Containerizing the Application on Docker:

1. Create a Dockerfile

```
touch Dockerfile
```

2. Create an Dockerfile based on TA's template

(<https://github.com/PrateekKumar1709/Docker-Demo/blob/main/cat-gif-app/single-container/Dockerfile>)

```

# Use an official Python runtime as the base image
FROM python:3.9-slim

# Set the working directory in the container
WORKDIR /app

# Copy the current directory contents into the container at /app
COPY . /app

# Upgrade pip and install the required packages
RUN pip install --no-cache-dir --upgrade pip && \
    pip install --no-cache-dir -r requirements.txt

# Make port 5050 available to the world outside this container
EXPOSE 5050

# Define environment variable
ENV NAME="World"

# Run app.py when the container launches
CMD ["python", "app.py"]

```

3. Build Docker

```
docker build -t flask-app .
```

```

yuwei@yuwei-SER:~/Documents/CC-A2/app$ docker images

```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
flask-app	latest	8dc1078b81aa	23 minutes ago	145MB

5. Write a docker compose (ref: https://hub.docker.com/_/mongo, <https://github.com/docker/compose>)

```

services:
  mongodb:
    image: "mongo:latest"
    ports:

```

- "27017:27017"

environment:

- MONGO_INITDB_DATABASE=camp2016

volumes:

- mongo_data:/data/db

flask_app:

image: "flask-app:latest" # Use pre-built image

ports:

- "5050:5050"

depends_on:

- mongodb

environment:

- MONGO_HOST=mongodb

- MONGO_PORT=27017

volumes:

mongo_data:

5. Start or stop services for the flask app and MongoDB container

docker compose up -d

docker compose down

```
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:11] "GET / HTTP/1.1" 200 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:11] "GET /static/assets/style.css HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:11] "GET /static/assets/emoji.js HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:11] "GET /static/assets/emoji.css HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:11] "GET /static/images/no.png HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:11] "GET /static/assets/twemoji.min.js HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:12] "GET /static/images/no.png HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:14] "GET /list HTTP/1.1" 200 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:14] "GET /static/assets/emoji.js HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:15] "GET /static/assets/style.css HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:15] "GET /static/images/yes.png HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:15] "GET /static/assets/emoji.css HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:15] "GET /static/assets/twemoji.min.js HTTP/1.1" 304 -
flask_app-1 | 172.80.1.1 - - [13/Mar/2025 01:31:15] "GET /static/images/no.png HTTP/1.1" 304 -
Gracefully stopping... (press Ctrl+C again to force)
[+] Stopping 2/2
  ✓ Container app-flask_app-1 Stopped 0.3s
  ✓ Container app-mongodb-1 Stopped 0.2s
yuwei@yuwei-SER:~/Documents/CC-A2/app$ docker compose down
[+] Running 3/3
  ✓ Container app-flask_app-1 Removed 0.0s
  ✓ Container app-mongodb-1 Removed 0.0s
  ✓ Network app_default Removed 0.2s
```

ToDo Reminder

ALL Uncompleted Completed About

Search Reference:

Unique ID ▼

Search Task

Search

To-Do LIST :

Status	Task Name	Description Name	Date	Priority	Remove	Modify
✓	Add_data	add some data	2025-03-08	Low !		
✗	Done	Completed Task	2025-03-07	Medium !!		

Add a Task

Taskname

Enter Description here...

mm/dd/yyyy 📅

Priority ▼

Create

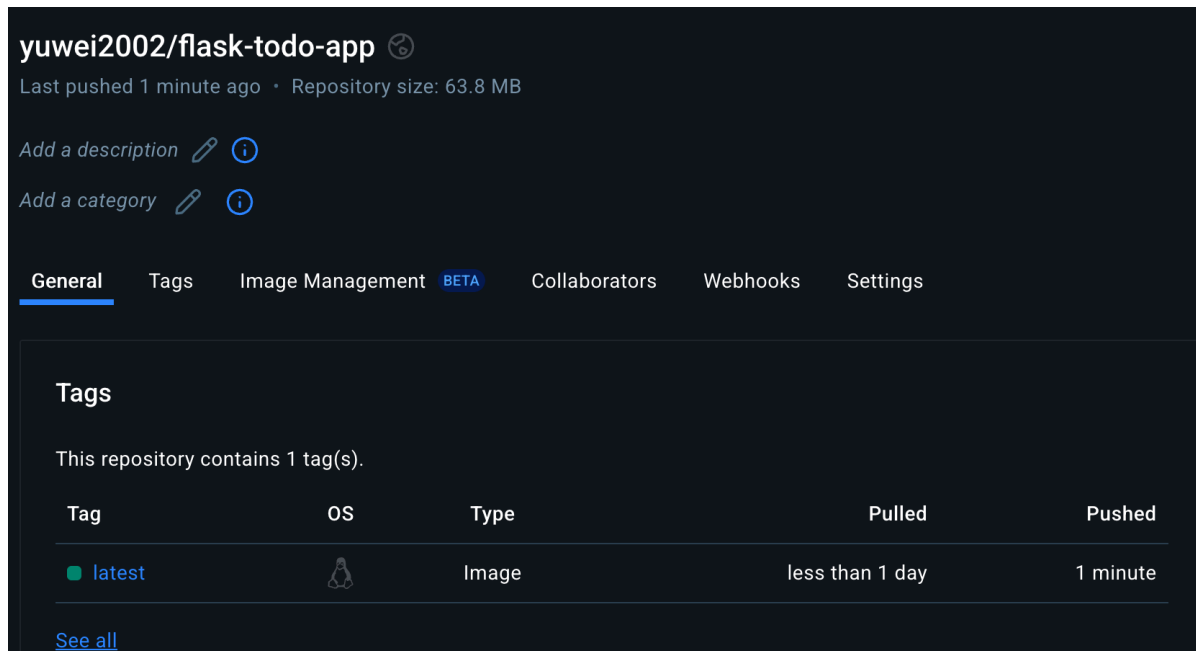
6. Tag that Docker image

```
docker tag app-flask_app yuwei2002/flask-todo-app:latest
```

7. Push the Docker image to Dockerhub

```
docker push yuwei2002/flask-todo-app:latest
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/app$ docker push yuwei2002/flask-todo-app:latest
The push refers to repository [docker.io/yuwei2002/flask-todo-app]
59c9ee8d6eba: Pushed
895c8499b912: Pushed
14eb7e6c0f1b: Layer already exists
70fff1cc441b: Layer already exists
62f5a620d172: Layer already exists
4a0e2c276a6a: Layer already exists
5f1ee22ffb5e: Layer already exists
latest: digest: sha256:ae7bce2b212fb56f0141cb428174a98887fff6e08b1c0d06e8cfff4ae4c07daba size: 1786
```



Part 3: Deploying the Application on Minikube:

1. Start Minikube using the command-line interface:

```
minikube start
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ minikube start
🐳 minikube v1.35.0 on Ubuntu 22.04
🔧 Using the docker driver based on existing profile
👉 Starting "minikube" primary control-plane node in "minikube" cluster
📦 Pulling base image v0.0.46 ...
🔄 Restarting existing docker container for "minikube" ...
🔧 Preparing Kubernetes v1.32.0 on Docker 27.4.1 ...
🔍 Verifying Kubernetes components...
   ▪ Using image gcr.io/k8s-minikube/storage-provisioner:v5
🌟 Enabled addons: storage-provisioner, default-storageclass
🏁 Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
```

2. Create two pods: one for the flask app and one for the MongoDB to store data.

```
# Flask-app
# Service
apiVersion: v1
kind: Service
```

```

metadata:
  name: flask-app
spec:
  selector:
    app: flask-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 5050
  type: LoadBalancer
---
# Deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask-app
spec:
  replicas: 1
  selector:
    matchLabels:
      app: flask-app
  template:
    metadata:
      labels:
        app: flask-app
    spec:
      containers:
        - name: flask-app
          image: yuwei2002/flask-todo-app:latest
          ports:
            - containerPort: 5050
          env:
            - name: MONGO_HOST
              value: "mongodb"
            - name: MONGO_PORT
              value: "27017"
          resources:

```

```
requests:
  cpu: "0.5"
  memory: "512Mi"
limits:
  memory: "512Mi"
  cpu: "1"
imagePullPolicy: Always
```

```
# MongoDB
# Persistent volume claim
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mongo-pvc
spec:
  accessModes:
    - ReadWriteOnce
  volumeMode: Filesystem
  storageClassName: standard # gp2 # AWS EBS storage
  resources:
    requests:
      storage: 1Gi
---
# Service
apiVersion: v1
kind: Service
metadata:
  name: mongodb
spec:
  selector:
    app: mongodb
  ports:
    - port: 27017
      targetPort: 27017
---
# Deployment
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mongodb
spec:
  selector:
    matchLabels:
      app: mongodb
  template:
    metadata:
      labels:
        app: mongodb
    spec:
      containers:
        - name: mongodb
          image: mongo:latest
          ports:
            - containerPort: 27017
          env:
            - name: MONGO_INITDB_DATABASE
              value: camp2016
          volumeMounts:
            - name: mongo-data
              mountPath: /data/db
      resources:
        requests:
          cpu: "0.3"
          memory: "512Mi"
        limits:
          memory: "512Mi"
          cpu: "0.5"
      volumes:
        - name: mongo-data
          persistentVolumeClaim:
            claimName: mongo-pvc
```

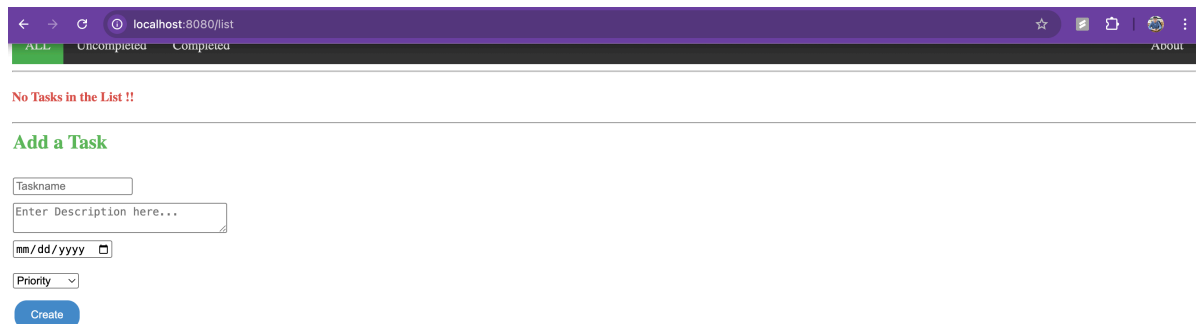
3. Apply these YAML using `kubectl`


```
kubectl apply -f app-cloud.yaml
kubectl apply -f mongo-cloud.yaml
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl apply -f app-cloud.yaml
service/flask-app unchanged
deployment.apps/flask-app created
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl apply -f mongo-cloud.yaml
persistentvolumeclaim/mongo-pvc unchanged
service/mongodb unchanged
deployment.apps/mongodb created
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get deployments
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
flask-app     1/1     1            1           9s
mongodb       1/1     1            1           6s
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-85cd75cc5-c27kt           1/1     Running   0          12s
mongodb-54bcd6bb4-vcwtb             1/1     Running   0          9s
yuwei@yuwei-SER:~/Documents/CC-A2/kube$
```

4. Test the application

```
# Port Forwarding
kubectl port-forward service/flask-app 8080:80 -n default
```



5. Cleanup and stop

```
kubectl delete deployments --all
minikube stop
```

Part 4: Deploying the Application on AWS EKS:

Ref: <https://docs.aws.amazon.com/eks/latest/userguide/getting-started.html>

1. Install AWS CLI

```
# Linux
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
unzip awscliv2.zip
sudo ./aws/install

#MacOS
brew install weaveworks/tap/eksctl
brew install awscli
```

2. Create credential file for cluster

```
aws configure
# Check identity
aws sts get-caller-identity
```

3. Setup eksctl

```
# for ARM systems, set ARCH to: `arm64`, `armv6` or `armv7`
ARCH=amd64
PLATFORM=$(uname -s)_$ARCH

curl -sLO "https://github.com/eksctl-io/eksctl/releases/latest/download/eksctl_${PLATFORM}.tar.gz"

# (Optional) Verify checksum
curl -sL "https://github.com/eksctl-io/eksctl/releases/latest/download/eksctl_checksums.txt"

tar -xzf eksctl_${PLATFORM}.tar.gz -C /tmp && rm eksctl_${PLATFORM}.tar.gz

sudo mv /tmp/eksctl /usr/local/bin
```

4. Create an AWS EKS cluster with an IAM using eksctl

```
# Create cluster
eksctl create cluster --name my-cluster --region us-east-2

# created IAM Open ID Connect provider
eksctl utils associate-iam-oidc-provider --region=us-east-2 --cluster=my-cluster
# Create IAM
eksctl create iamserviceaccount \
    --name ebs-csi-controller-sa \
    --namespace kube-system \
    --cluster my-cluster \
    --role-name AmazonEKS_EBS_CSI_DriverRole \
    --role-only \
    --attach-policy-arn arn:aws:iam::aws:policy/service-role/AmazonEBSCSIDriverPolicy \
    --approve
```

5. Install `aws-ebs-csi-driver` since we using AWS EBS storage

```
eksctl create addon --name aws-ebs-csi-driver --cluster my-cluster --region us-east-2
```

6. Configure the Kubernetes CLI (kubectl) to connect to the EKS cluster

Configure the Kubernetes CLI (kubectl) to connect to the EKS cluster

Check connection

```
kubectl get nodes -o wide
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE   VERSION            INTERNAL-IP    EXTERNAL-IP
ip-192-168-3-45.us-east-2.compute.internal Ready    <none>   3d23h v1.30.9-eks-5d632ec 192.168.3.45   3.145.189.5
Amazon Linux 2 5.10.234-225.895.amzn2.x86_64 containerd://1.7.25
ip-192-168-48-154.us-east-2.compute.internal Ready    <none>   3d23h v1.30.9-eks-5d632ec 192.168.48.154 3.137.142.241
Amazon Linux 2 5.10.234-225.895.amzn2.x86_64 containerd://1.7.25
```

7. Deployment (based on previous YAML in Part 3)

Step 1: Change the `storageClassName` in mongo-cloud.yaml from **standard** to **gp2** since we want to use the AWS EBS storage

Step 2: Create a namespace for management

```
kubectl create namespace eks-todo-list-app
```

Step 3: Apply yaml

```
kubectl apply -f mongo-cloud.yaml -n eks-todo-list-app
```

```
kubectl apply -f app-cloud.yaml -n eks-todo-list-app
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl apply -f mongo-cloud.yaml -n eks-todo-list-app
persistentvolumeclaim/mongo-pvc created
service/mongodb created
deployment.apps/mongodb created
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl apply -f app-cloud.yaml -n eks-todo-list-app
service/flask-app created
deployment.apps/flask-app created
```

Check If successful:

```
kubectl get pods -n eks-todo-list-app
```

```
kubectl get all -n eks-todo-list-app
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get pods -n eks-todo-list-app
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-755bc58598-77b2c         1/1     Running   0           34s
mongodb-6f84848985-6hf85           1/1     Running   0           38s
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get all -n eks-todo-list-app
NAME                                READY   STATUS    RESTARTS   AGE
pod/flask-app-755bc58598-77b2c      1/1     Running   0           38s
pod/mongodb-6f84848985-6hf85        1/1     Running   0           42s

NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP                                     PORT(S)
service/flask-app                    LoadBalancer        10.100.105.37    adbe2494cccbf4be29d908a3e871e02b-139878156.us-east-2.elb.amazonaws.com    80
service/mongodb                      ClusterIP            10.100.196.43    <none>                                         27

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/flask-app            1/1     1             1           38s
deployment.apps/mongodb              1/1     1             1           43s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/flask-app-755bc58598 1          1         1       38s
replicaset.apps/mongodb-6f84848985    1          1         1       43s
yuwei@yuwei-SER:~/Documents/CC-A2/kube$
```

8. Test the application

```
kubectl get service flask-app -n eks-todo-list-app
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get service flask-app -n eks-todo-list-app
NAME    TYPE                CLUSTER-IP      EXTERNAL-IP                                     PORT(S)
flask-app LoadBalancer        10.100.105.37    adbe2494cccbf4be29d908a3e871e02b-139878156.us-east-2.elb.amazonaws.com    80:32566/TCP
CP      2m33s
```

The screenshot shows a web browser window with a purple header bar. Below the header, a red message states "No Tasks in the List !!". Underneath this is a green heading "Add a Task". The form contains several input fields: "Taskname", "Enter Description here..." (with a text area icon), "mm/dd/yyyy" (with a calendar icon), and a "Priority" dropdown menu. At the bottom of the form is a blue "Create" button.

Part 5: Replication controller feature:

1. Create an app-cloud.yaml to use ReplicationController instead of Deployment

```
# ReplicationController
apiVersion: v1
kind: ReplicationController
metadata:
  name: flask-app-rc
  namespace: eks-todo-list-app
spec:
  replicas: 5
  selector:
    app: flask-app-rc
  template:
    metadata:
      labels:
        app: flask-app-rc
    spec:
      containers:
        - name: flask-app
          image: yuwei2002/flask-todo-app:latest
          ports:
            - containerPort: 5050
          env:
            - name: MONGO_HOST
              value: "mongodb"
            - name: MONGO_PORT
```

```
    value: "27017"
  resources:
    requests:
      cpu: "0.3"
      memory: "256Mi"
    limits:
      cpu: "0.5"
      memory: "512Mi"
```

2. Create the replication controller :

```
kubectl apply -f app-rc.yaml -n eks-todo-list-app
```

verify that the specified number of replicas are created and running

```
kubectl get rc -n eks-todo-list-app
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl apply -f app-rc.yaml -n eks-todo-list-app
replicationcontroller/flask-app-rc created
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get rc -n eks-todo-list-app
NAME          DESIRED   CURRENT   READY   AGE
flask-app-rc   5         5         3       13s
```

3. Test the replication controller:

Step 1: Get pods name

```
kubectl get pods -l app=flask-app-rc -n eks-todo-list-app
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get pods -l app=flask-app-rc -n eks-todo-list-app
NAME                READY   STATUS    RESTARTS   AGE
flask-app-rc-8bkwm  1/1     Running   0          12s
flask-app-rc-cdjhn  1/1     Running   0          12s
flask-app-rc-dcptc  1/1     Running   0          12s
flask-app-rc-l678n  1/1     Running   0          12s
flask-app-rc-mqldv  1/1     Running   0          12s
```

Step 2: Delete a pod (flask-app-rc-52tzx)

```
kubectl delete pod flask-app-rc-8bkwm -n eks-todo-list-app
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl delete pod flask-app-rc-8bkwm -n eks-todo-list-app
pod "flask-app-rc-8bkwm" deleted
```

Step 3: Watch if a new pod auto running:

```
kubectl get pods -l app=flask-app-rc -n eks-todo-list-app -w
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get pods -l app=flask-app-rc -n eks-todo-list-app -w
NAME                READY    STATUS    RESTARTS   AGE
flask-app-rc-cdjhn   1/1      Running   0           65s
flask-app-rc-dcptc   1/1      Running   0           65s
flask-app-rc-jlmtf   1/1      Running   0           24s
flask-app-rc-l678n   1/1      Running   0           65s
flask-app-rc-mqldv   1/1      Running   0           65s
```

4. Update number of replicas and verify

Step 1: Change `replicas` from **5** to **3**

Step 2: Apply YAML:

```
kubectl apply -f app-rc.yaml -n eks-todo-list-app
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl apply -f app-rc.yaml -n eks-todo-list-app
replicationcontroller/flask-app-rc configured
```

Step 3: Verify:

```
kubectl get rc -n eks-todo-list-app
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get rc -n eks-todo-list-app
NAME          DESIRED    CURRENT    READY    AGE
flask-app-rc  3          3          3        3m4s
```

Part 6: Rolling update strategy:

1. Since replication controller do not support update strategy, we need back to Deployment

```
# Flask-app
# Service
```

```
apiVersion: v1
kind: Service
metadata:
  name: flask-app
spec:
  selector:
    app: flask-app-rc
  ports:
    - protocol: TCP
      port: 80
      targetPort: 5050
  type: LoadBalancer
---
# Deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask-app-deployment
  namespace: eks-todo-list-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: flask-app-rc
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  template:
    metadata:
      labels:
        app: flask-app-rc
    spec:
      containers:
        - name: flask-app
          image: yuwei2002/flask-todo-app:v2 # latest
```




```
ports:
  - containerPort: 5050
env:
  - name: MONGO_HOST
    value: "mongodb"
  - name: MONGO_PORT
    value: "27017"
resources:
  requests:
    cpu: "0.3"
    memory: "256Mi"
  limits:
    cpu: "0.5"
    memory: "512Mi"
```

2. Apply YAML



```
kubectl apply -f app-cloud.yaml -n eks-todo-list-app
```



3. Update the Docker image for the deployment to a new version

```
docker build -t yuwei2002/flask-todo-app:v2 .
docker push yuwei2002/flask-todo-app:v2
```

yuwei2002/flask-todo-app 

Last pushed less than a minute ago · Repository size: 70.3 MB





Add a description  

Add a category  

General Tags Image Management BETA Collaborators Webhooks Settings

Tags

This repository contains 2 tag(s).

Tag	OS	Type	Pulled	Pushed
 v2		Image	less than 1 day	less than a minute
 latest		Image	less than 1 day	about 2 hours

[See all](#)

4. Update YAML to use new docker: Change tag in image

5. Apply YAML

```
kubectl apply -f kube/app-cloud.yaml -n eks-todo-list-app
```

6. Monitor the rolling update progress

```
kubectl rollout status deployment/flask-app-deployment -n eks-todo-list-app
# Watch pods being updated
kubectl get pods -n eks-todo-list-app -l app=flask-app-rc -w
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl rollout status deployment/flask-app-deployment -n eks-todo-list-app
deployment "flask-app-deployment" successfully rolled out
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get pods -n eks-todo-list-app -l app=flask-app-rc -w
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-588d47d8c6-48f77 1/1     Running   0           2m15s
flask-app-deployment-588d47d8c6-f5vzl 1/1     Running   0           2m15s
flask-app-deployment-588d47d8c6-f9mxf 1/1     Running   0           2m15s
```

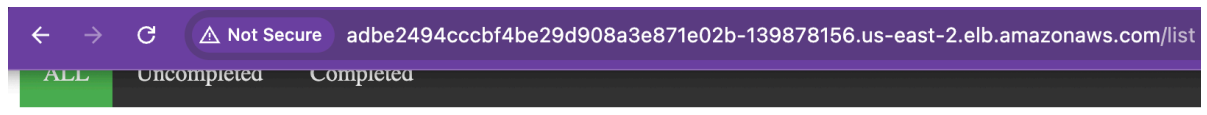
7. Test the application:

```
kubectl get service flask-app -n eks-todo-list-app
```

```

yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get service flask-app -n eks-todo-list-app
NAME      AGE      TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)
flask-app 57m      LoadBalancer  10.100.105.37    adbe2494cccbf4be29d908a3e871e02b-139878156.us-east-2.elb.amazonaws.com  80:32566/T

```



No Tasks in the List !!

Add a Task

7. Ensure application using New Docker image version

```

yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl describe deployment flask-app-deployment -n eks-todo-list-app
Name: flask-app-deployment
Namespace: eks-todo-list-app
CreationTimestamp: Wed, 12 Mar 2025 23:29:25 -0400
Labels: <none>
Annotations: deployment.kubernetes.io/revision: 1
Selector: app=flask-app-rc
Replicas: 3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 1 max unavailable, 1 max surge
Pod Template:
  Labels: app=flask-app-rc
  Containers:
    flask-app:
      Image: yuwei2002/flask-todo-app:v2
      Port: 5050/TCP
      Host Port: 0/TCP
      Limits:
        cpu: 500m
        memory: 512Mi
      Requests:
        cpu: 300m
        memory: 256Mi
      Environment:
        MONGO_HOST: mongodb
        MONGO_PORT: 27017
      Mounts: <none>
      Volumes: <none>
      Node-Selectors: <none>
      Tolerations: <none>
  Conditions:
    Type           Status    Reason
    ----           -
    Available       True      MinimumReplicasAvailable
    Progressing     True      NewReplicaSetAvailable

```

Part 7: Health monitoring:

1. Add health and ready for probe in `app.py`.

```

from flask import Flask, render_template, request, redirect, url_for # For flask implementation
from pymongo import MongoClient # Database connector
from bson.objectid import ObjectId # For ObjectId to work
from bson.errors import InvalidId # For catching InvalidId exception for ObjectId
import os
import threading # Add threading for timer functionality

mongodb_host = os.environ.get('MONGO_HOST', 'localhost')
mongodb_port = int(os.environ.get('MONGO_PORT', '27017'))
client = MongoClient(mongodb_host, mongodb_port) #Configure the connection

```

```

db = client.camp2016 #Select the database
todos = db.todo #Select the collection

app = Flask(__name__)
title = "TODO with Flask(V2)"
heading = "ToDo Reminder"
#modify=ObjectId()

# Global variables to control health and readiness status
app_healthy = True
app_ready = True
ready_timer = None # Timer to track auto-reset

def redirect_url():
    return request.args.get('next') or \
        request.referrer or \
        url_for('index')

@app.route("/list")
def lists ():
    #Display the all Tasks
    todos_l = todos.find()
    a1="active"
    return render_template('index.html',a1=a1,todos=todos_l,t=title,h=heading)

@app.route("/")
@app.route("/uncompleted")
def tasks ():
    #Display the Uncompleted Tasks
    todos_l = todos.find({"done":"no"})
    a2="active"
    return render_template('index.html',a2=a2,todos=todos_l,t=title,h=heading)

@app.route("/completed")
def completed ():
    #Display the Completed Tasks

```

```

    todos_l = todos.find({"done":"yes"})
    a3="active"
    return render_template('index.html',a3=a3,todos=todos_l,t=title,h=heading)

@app.route("/done")
def done ():
    #Done-or-not ICON
    id=request.values.get("_id")
    task=todos.find({"_id":ObjectId(id)})
    if(task[0]["done"]=="yes"):
        todos.update_one({"_id":ObjectId(id)}, {"$set": {"done":"no"}})
    else:
        todos.update_one({"_id":ObjectId(id)}, {"$set": {"done":"yes"}})
    redir=redirect_url() # Re-directed URL i.e. PREVIOUS URL from where it cam

    # if(str(redir)=="http://localhost:5000/search"):
    #     redir+="?key="+id+"&refer="+refer

    return redirect(redir)

#@app.route("/add")
#def add():
# return render_template('add.html',h=heading,t=title)

@app.route("/action", methods=['POST'])
def action ():
    #Adding a Task
    name=request.values.get("name")
    desc=request.values.get("desc")
    date=request.values.get("date")
    pr=request.values.get("pr")
    todos.insert_one({ "name":name, "desc":desc, "date":date, "pr":pr, "done":"no" })
    return redirect("/list")

@app.route("/remove")
def remove ():
    #Deleting a Task with various references

```

```

key=request.values.get("_id")
todos.delete_one({"_id":ObjectId(key)})
return redirect("/")

@app.route("/update")
def update ():
    id=request.values.get("_id")
    task=todos.find({"_id":ObjectId(id)})
    return render_template('update.html',tasks=task,h=heading,t=title)

@app.route("/action3", methods=['POST'])
def action3 ():
    #Updating a Task with various references
    name=request.values.get("name")
    desc=request.values.get("desc")
    date=request.values.get("date")
    pr=request.values.get("pr")
    id=request.values.get("_id")
    todos.update_one({"_id":ObjectId(id)}, {'$set':{ "name":name, "desc":desc, "d
    return redirect("/")

@app.route("/search", methods=['GET'])
def search():
    #Searching a Task with various references

    key=request.values.get("key")
    refer=request.values.get("refer")
    if(refer=="id"):
        try:
            todos_l = todos.find({refer:ObjectId(key)})
            if not todos_l:
                return render_template('index.html',a2=a2,todos=todos_l,t=title,h=heac
        except InvalidId as err:
            pass
            return render_template('index.html',a2=a2,todos=todos_l,t=title,h=headin
    else:
        todos_l = todos.find({refer:key})

```

```

    return render_template('searchlist.html', todos=todos_l, t=title, h=heading)

@app.route("/about")
def about():
    return render_template('credits.html', t=title, h=heading)

# livenessProbe
@app.route("/health")
def health():
    if app_healthy:
        return {"status": "healthy"}, 200
    else:
        return {"status": "unhealthy"}, 500

# readinessProbe
@app.route("/ready")
def ready():
    if not app_ready:
        return {"status": "not ready", "error": "Readiness manually disabled"}, 503
    try:
        client.server_info() # Check if MongoDB connection is still alive
        return {"status": "ready"}, 200
    except Exception as e:
        return {"status": "not ready", "error": str(e)}, 503

# reset timer for /ready
def reset_ready():
    global app_ready, ready_timer
    app_ready = True
    ready_timer = None
    print("Readiness automatically reset to True after 30 seconds")

# Test endpoints to simulate failures
@app.route("/toggle-health")
def toggle_health():
    global app_healthy
    app_healthy = not app_healthy
    status = "unhealthy" if not app_healthy else "healthy"

```



```

    return {"message": f"Health status toggled to {status}"}, 200

# Test endpoint to toggle readiness status
@app.route("/toggle-ready")
def toggle_ready():
    global app_ready, ready_timer
    app_ready = not app_ready
    status = "not ready" if not app_ready else "ready"
    # If toggled to not ready, set timer to reset after 30 seconds
    if not app_ready:
        if ready_timer:
            ready_timer.cancel()
        ready_timer = threading.Timer(30.0, reset_ready)
        ready_timer.daemon = True # Make thread daemon so it won't block app sh
        ready_timer.start()
        return {"message": f"Readiness status toggled to {status}. Will reset in 30 s
    else:
        if ready_timer:
            ready_timer.cancel()
            ready_timer = None
        return {"message": f"Readiness status toggled to {status}"}, 200

if __name__ == "__main__":
    env = os.environ.get('FLASK_ENV', 'development')
    port = int(os.environ.get('PORT', 5050))
    debug = False if env == 'production' else True
    app.run(host='0.0.0.0', port=port, debug=debug)
    # Careful with the debug mode..

```

2. Build a new docker for app.py and push to Dockerhub

```

docker build -t yuwei2002/flask-todo-app:v3 .
docker push yuwei2002/flask-todo-app:v3

```

```

yuwei@yuwei-SER:~/Documents/CC-A2/app$ docker build -t yuwei2002/flask-todo-app:v3 .
[+] Building 4.1s (10/10) FINISHED                                docker:default
=> [internal] load build definition from Dockerfile                0.0s
=> => transferring dockerfile: 588B                               0.0s
=> [internal] load metadata for docker.io/library/python:3.9-slim 0.2s
=> [auth] library/python:pull token for registry-1.docker.io      0.0s
=> [internal] load .dockerignore                                  0.0s
=> => transferring context: 2B                                     0.0s
=> [1/4] FROM docker.io/library/python:3.9-slim@sha256:d1fd807555208707ec95b2 0.0s
=> [internal] load build context                                  0.0s
=> => transferring context: 6.84kB                                 0.0s
=> CACHED [2/4] WORKDIR /app                                       0.0s
=> [3/4] COPY . /app                                              0.0s
=> [4/4] RUN pip install --no-cache-dir --upgrade pip &&         pip install --no 3.8s
=> exporting to image                                              0.1s
=> => exporting layers                                           0.1s
=> => writing image sha256:67a9c2b269053128a0f5a0bb640bac2c7846db72525ce2e742 0.0s
=> => naming to docker.io/yuwei2002/flask-todo-app:v3           0.0s

```

```

yuwei@yuwei-SER:~/Documents/CC-A2/app$ docker push yuwei2002/flask-todo-app:v3
The push refers to repository [docker.io/yuwei2002/flask-todo-app]
ce1de7ed81ec: Pushed
9f97b91bbb57: Pushed
14eb7e6c0f1b: Layer already exists
70fff1cc441b: Layer already exists
62f5a620d172: Layer already exists
4a0e2c276a6a: Layer already exists
5f1ee22ffb5e: Layer already exists
v3: digest: sha256:2f87d1b38fa00cdd23434ef6347d9e7cede759582e479f1b82957fb955ab4825 size: 1786

```

yuwei2002/flask-todo-app

Last pushed less than a minute ago · Repository size: 89.6 MB

Add a description
Add a category

General
Tags
Image Management
Collaborators
Webhooks
Settings

Tags

This repository contains 3 tag(s).

Tag	OS	Type	Pulled	Pushed
v3	linux	Image	less than 1 day	less than a minute
v2	linux	Image	less than 1 day	about 17 hours
latest	linux	Image	less than 1 day	about 19 hours

3. Add livenessProbe and readinessProbe to app-cloud.yaml

```
# Flask-app
# Service
apiVersion: v1
kind: Service
metadata:
  name: flask-app
spec:
  selector:
    app: flask-app-rc
  ports:
    - protocol: TCP
      port: 80
      targetPort: 5050
  type: LoadBalancer
---
# Deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask-app-deployment
  namespace: eks-todo-list-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: flask-app-rc
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  template:
    metadata:
      labels:
        app: flask-app-rc
    spec:
```

```
containers:
  - name: flask-app
    image: yuwei2002/flask-todo-app:v3 # latest
    ports:
      - containerPort: 5050
    env:
      - name: MONGO_HOST
        value: "mongodb"
      - name: MONGO_PORT
        value: "27017"
    resources:
      requests:
        cpu: "0.3"
        memory: "256Mi"
      limits:
        cpu: "0.5"
        memory: "512Mi"
    imagePullPolicy: Always
    livenessProbe: # Http Liveness Probe
      httpGet:
        path: /health
        port: 5050
      initialDelaySeconds: 30
      periodSeconds: 10
      timeoutSeconds: 5
      failureThreshold: 3
    readinessProbe: # Http Readiness Probe
      httpGet:
        path: /ready
        port: 5050
      initialDelaySeconds: 15
      periodSeconds: 5
      timeoutSeconds: 3
      successThreshold: 1
      failureThreshold: 3
```

4. Apply the updated YAML TODO

```
kubectl apply -f app-cloud-copy.yaml -n eks-todo-list-app
```

```
Cyuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get pods -n eks-todo-list-app
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-68d8f9f787-9d744 1/1     Running   0           89s
flask-app-deployment-68d8f9f787-ltwg5 1/1     Running   0           89s
flask-app-deployment-68d8f9f787-pbtpl 1/1     Running   0           89s
mongodb-6f84848985-bc8rz              1/1     Running   0          5m17s
```

5. Get external IP of deployment

```
kubectl get service flask-app -n eks-todo-list-app
```

```
# a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get service flask-app -n eks-todo-list-app
NAME    TYPE        CLUSTER-IP    EXTERNAL-IP                                PORT(S)
flask-app LoadBalancer 10.100.98.145  a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com 80:31128/TCP
CP      5m18s
```

6. Monitor health of a pod (using pod name)

```
kubectl describe pod flask-app-deployment-68d8f9f787-9d744 -n eks-todo-list
```

```
node.kubernetes.io/unreachable:NoExecute op=Exists for 300s

Events:
  Type     Reason      Age   From              Message
  ----     ------      ---   -
  Normal   Scheduled   10m   default-scheduler Successfully assigned eks-todo-list-app/flask-app-deployment-68d8f9f787-9d744 to ip-192-168-36-118.us-east-2.compute.internal
  Normal   Pulled      10m   kubelet           Successfully pulled image "yuwei2002/flask-todo-app:v3" in 1.503s (1.503s including waiting). Image size: 53425846 bytes.
  Normal   Pulling     102s (x2 over 10m)  kubelet           Pulling image "yuwei2002/flask-todo-app:v3"
  Normal   Created     102s (x2 over 10m)  kubelet           Created container flask-app
  Normal   Started     102s (x2 over 10m)  kubelet           Started container flask-app
```

7. Trigger an active readiness prob via `curl`

```
curl http://a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/kube$ curl http://a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com/toggle-ready
{"message": "Readiness status toggled to not ready. Will reset in 30 seconds."}
```

To better test the liveness prob, we set the duration of "not ready" to 30s.

```

yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get pods -n eks-todo-list-app -w
NAME                                READY    STATUS    RESTARTS   AGE
flask-app-deployment-68d8f9f787-9d744 1/1      Running   0           4m44s
flask-app-deployment-68d8f9f787-ltwg5 1/1      Running   0           4m44s
flask-app-deployment-68d8f9f787-pbppl 1/1      Running   0           4m44s
mongodb-6f84848985-bc8rz              1/1      Running   0           8m32s
flask-app-deployment-68d8f9f787-ltwg5 0/1      Running   0           4m46s
flask-app-deployment-68d8f9f787-ltwg5 1/1      Running   0           5m6s

```

8. Trigger an active readiness prob via `curl`

```
curl http://a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazo
```

```

yuwei@yuwei-SER:~/Documents/CC-A2/kube$ curl http://a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com/toggle-health
{
  "message": "Health status toggled to unhealthy"
}

```

Once liveness prob detect failure, it will restart the pod.

```

yuwei@yuwei-SER:~/Documents/CC-A2/kube$ kubectl get pods -n eks-todo-list-app -w
NAME                                READY    STATUS    RESTARTS   AGE
flask-app-deployment-68d8f9f787-9d744 1/1      Running   0           8m54s
flask-app-deployment-68d8f9f787-ltwg5 1/1      Running   0           8m54s
flask-app-deployment-68d8f9f787-pbppl 1/1      Running   0           8m54s
mongodb-6f84848985-bc8rz              1/1      Running   0           12m
flask-app-deployment-68d8f9f787-9d744 0/1      Running   1 (1s ago)  9m12s
flask-app-deployment-68d8f9f787-9d744 1/1      Running   1 (15s ago) 9m26s

```

Result of health monitor:

```

node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age    From          Message
  ----    -
  Normal  Scheduled   10m    default-scheduler  Successfully assigned eks-todo-list-app/flask-app-deployment-68d8f9f787-9d744 to ip-192-168-36-118.us-east-2.compute.internal
  Normal  Pulled      10m    kubelet        Successfully pulled image "yuwei2002/flask-todo-app:v3" in 1.503s (1.503s including waiting). Image size: 53425846 bytes.
  Normal  Pulling     102s (x2 over 10m)  kubelet        Pulling image "yuwei2002/flask-todo-app:v3"
  Normal  Created     102s (x2 over 10m)  kubelet        Created container flask-app
  Normal  Started     102s (x2 over 10m)  kubelet        Started container flask-app
  Warning  Unhealthy   102s (x3 over 2m2s) kubelet        Liveness probe failed: HTTP probe failed with statuscode: 500
  Normal  Killing     102s    kubelet        Container flask-app failed liveness probe, will be restarted
  Normal  Pulled      102s    kubelet        Successfully pulled image "yuwei2002/flask-todo-app:v3" in 173ms (173ms including waiting). Image size: 53425846 bytes.

```

Step 8: Alerting:

1. Install Prometheus

```
wget https://github.com/prometheus/prometheus/releases/download/v3.2.1/prometheus-3.2.1.linux-amd64.tar.gz
```

```
tar xvfz prometheus-*.tar.gz
```

```
cd prometheus-3.2.1.linux-amd64/
```

2. Configure Prometheus

Step 1: modify prometheus.yml

```
global:
  scrape_interval: 15s # collect metrics from targets every 15 seconds
  evaluation_interval: 15s

rule_files:
  - "probe-alerts.yml"

alerting:
  alertmanagers:
    - static_configs:
        - targets:
            - alertmanager:9093

scrape_configs:
  - job_name: "all-kubernetes-pods"
    kubernetes_sd_configs:
      - role: pod
    relabel_configs:
      # Add namespace as a label
      - source_labels: [__meta_kubernetes_namespace]
        action: replace
        target_label: kubernetes_namespace
      # Add pod name as a label
      - source_labels: [__meta_kubernetes_pod_name]
        action: replace
        target_label: kubernetes_pod_name

    # Uses direct HTTP checks instead
    - job_name: "blackbox"
      metrics_path: /metrics
      static_configs:
```

```

- targets:
  - localhost:9115

# Direct health check
- job_name: "app-health-check"
  metrics_path: /health
  scrape_interval: 10s
  static_configs:
    - targets:
      - a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com
  relabel_configs:
    - source_labels: [__address__]
      target_label: instance
    - target_label: __address__
      replacement: a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com

```

step 2: Create probe-alerts.yml

```
touch probe-alerts.yml
```

```

groups:
- name: probe-alerts
  rules:
    - alert: ProbeFailureThresholdExceeded
      expr: probe_success == 0
      for: 45s
      labels:
        severity: critical
      annotations:
        summary: "Probe failure detected"
        description: "Probe {{ $labels.instance }} has been failing for 45 seconds,"

    - alert: HighProbeFailureRate
      expr: sum(rate(probe_success{job="blackbox"}[5m]) == 0) / sum(rate(probe_success{job="blackbox"}[5m]) > 0) > 0.5
      for: 45s
      labels:

```



```
severity: warning
annotations:
  summary: "High probe failure rate detected"
  description: "More than 10% of probes are failing in the last 10 minutes."

- alert: SlowResponseTime
  expr: probe_duration_seconds > 1
  for: 45s
  labels:
    severity: warning
  annotations:
    summary: "Slow probe response time"
    description: "Probe {{ $labels.instance }} response time is above 1 second"
```

3. Create a webhooks

Step 1: Following <https://api.slack.com/messaging/webhooks> to create a webhooks

Step 2: Test if success

```
curl -X POST -H 'Content-type: application/json' --data '{"text":"Hello, World!"}'
```

4. Install Alertmanager

```
wget https://github.com/prometheus/alertmanager/releases/download/v0.28.1/alertmanager-0.28.1.linux-amd64.tar.gz
tar xvfz alertmanager-*.tar.gz
cd alertmanager-0.28.1.linux-amd64/
```

5. Configure Alertmanager

Step 1: Modify alertmanager.yml

```
global:
  resolve_timeout: 1m # Faster resolution notifications
  slack_api_url: "https://hooks.slack.com/services/T08AJ5GLFKM/B08J048A2P6"
```

```

route:
  group_by: ["alertname", "job"]
  group_wait: 10s
  group_interval: 45s
  repeat_interval: 4h
  receiver: "slack-notifications"

receivers:
  - name: "slack-notifications"
    slack_configs:
      - channel: "#alerts"
        send_resolved: true
        title: "{{ range .Alerts }}{{ .Annotations.summary }}\n{{ end }}"
        text: "{{ range .Alerts }}{{ .Annotations.description }}\n{{ end }}"
        color: '{{ if eq .Status "firing" }}danger{{ else }}good{{ end }}'

```

6. Testing the Alerting System

Step 1: Start alertmanager

```
./alertmanager --config.file=alertmanager.yml &
```

```

yuwei@yuwei-SER:~/Documents/CC-A2/alertmanager-0.28.1.linux-amd64$ ./alertmanager --config
g.file=alertmanager.yml &
[1]+  Done                  ./alertmanager --config.file=alertmanager.yml
[1] 1382593
yuwei@yuwei-SER:~/Documents/CC-A2/alertmanager-0.28.1.linux-amd64$ time=2025-03-15T15:45:
19.485Z level=INFO source=main.go:191 msg="Starting Alertmanager" version="(version=0.28.
1, branch=HEAD, revision=b2099eaa2c9ebc25edb26517cb9c732738e93910)"
time=2025-03-15T15:45:19.485Z level=INFO source=main.go:192 msg="Build context" build_con
text="(go=go1.23.7, platform=linux/amd64, user=root@fa3ca569dfe4, date=20250307-15:05:18,
tags=netgo)"
time=2025-03-15T15:45:19.487Z level=INFO source=cluster.go:185 msg="setting advertise add
ress explicitly" component=cluster addr=10.12.66.12 port=9094
time=2025-03-15T15:45:19.488Z level=INFO source=cluster.go:674 msg="Waiting for gossip to
settle..." component=cluster interval=2s
time=2025-03-15T15:45:19.502Z level=INFO source=coordinator.go:112 msg="Loading configura
tion file" component=configuration file=alertmanager.yml
time=2025-03-15T15:45:19.503Z level=INFO source=coordinator.go:125 msg="Completed loading
of configuration file" component=configuration file=alertmanager.yml
time=2025-03-15T15:45:19.504Z level=INFO source=tls_config.go:347 msg="Listening on" addr
ess=[::]:9093
time=2025-03-15T15:45:19.504Z level=INFO source=tls_config.go:350 msg="TLS is disabled."
http2=false address=[::]:9093
time=2025-03-15T15:45:21.488Z level=INFO source=cluster.go:699 msg="gossip not settled" c
omponent=cluster polls=0 before=0 now=1 elapsed=2.000077995s
time=2025-03-15T15:45:29.490Z level=INFO source=cluster.go:691 msg="gossip settled; proce
ding" component=cluster elapsed=10.002617714s
yuwei@yuwei-SER:~/Documents/CC-A2/alertmanager-0.28.1.linux-amd64$ 

```

Step 2: Start prometheus

```
cd ../prometheus-3.2.1.linux-amd64/

./prometheus --config.file=prometheus.yml &
```

```
time=2025-03-15T15:46:05.623Z level=INFO source=main.go:1575 msg="See you next time."
yuwei@yuwei-SER:~/Documents/CC-A2/prometheus-3.2.1.linux-amd64$ ./prometheus --config.file=prometheus.yml &
[1]+  Done                  ./prometheus --config.file=prometheus.yml
[1] 1383081
yuwei@yuwei-SER:~/Documents/CC-A2/prometheus-3.2.1.linux-amd64$ time=2025-03-15T15:46:05.623Z level=INFO source=main.go:640 msg="No time or size retention was set so using the default time retention" duration=15d
time=2025-03-15T15:46:05.623Z level=INFO source=main.go:687 msg="Starting Prometheus Server" mode=server version="(version=3.2.1, branch=HEAD, revision=804c49d58f3f3784c77c9c8ec17c9062092cae27)"

time=2025-03-15T15:46:05.632Z level=INFO source=main.go:1486 msg="Completed loading of configuration file" db_storage=930ns remote_storage=1.08µs web_handler=200ns query_engine=600ns scrape=320.448µs scrape_sd=46.15µs notify=77.05µs notify_sd=10.11µs rules=485.587µs tracing=2.92µs filename=prometheus.yml totalDuration=1.377512ms
time=2025-03-15T15:46:05.632Z level=INFO source=main.go:1213 msg="Server is ready to receive web requests."
time=2025-03-15T15:46:05.632Z level=INFO source=manager.go:175 msg="Starting rule manager..." component="rule manager"
yuwei@yuwei-SER:~/Documents/CC-A2/prometheus-3.2.1.linux-amd64$
```

Step 3: Trigger failure

```
curl http://a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amaz
```

```
yuwei@yuwei-SER:~/Documents/CC-A2/prometheus-3.2.1.linux-amd64$ curl http://a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com/toggle-health
{
  "message": "Health status toggled to unhealthy"
}
```

Step 4: Watch firing:

Open the webui (<http://localhost:9090>)

Prometheus Query Alerts Status

Filter by rule state Filter by rule name or labels

probe-alerts probe-alerts.yml **FIRING (2)**

AppHealthFailureDetected **FIRING (1)**

up{job="app-health-check"} == 0 or http_status_code != 200

for: 45s

severity="critical"

description The health check for {{ \$labels.instance }} has been failing for 45 seconds, which may trigger Kubernetes restarts.

summary Application health check failing

Alert labels

alertnames	instances	job	State	Active Since	Value
AppHealthFailureDetected	a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com	app-health-check	FIRING	2m 44.504s	0

description The health check for a178b2df693a147af84fb7579c29c31a-115792124.us-east-2.elb.amazonaws.com has been failing for 45 seconds, which may trigger Kubernetes restarts.

summary Application health check failing

Slack Update:

Prometheus-alert APP 10:35 AM
Hello, World!

Prometheus-alert APP 12:23 PM

Pod flask-app-deployment-58d4d89d59-2xlfr in namespace eks-todo-list-app is failing liveness probes

Pod flask-app-deployment-58d4d89d59-2xlfr in namespace eks-todo-list-app has been failing liveness probes for more than 5 minutes.

Prometheus | KubePodFailingLivenessProbe